JETIR.ORG ISSN: 2349-5162 | ESTD Year : 2014 | Monthly Issue
JOURNAL OF EMERGING TECHNOLOGIES AND INNOVATIVE RESEARCH (JETIR)
An International Scholarly Open Access, Peer-reviewed, Refereed Journal

Opportunities and Challenges of Agentic AI in Finance

[1]Saaniya Chugh, [2]Aditya Vilas Deshpande

[1]Senior Technical Consultant, [2]Senior Software Engineer
[1]Global Customer Success,
[1]ServiceNow, Montreal, Canada

Abstract: The integration of Agentic AI in the finance domain is revolutionizing the way institutions handle fraud detection, portfolio management, risk assessment, and automated trading. The traditional finance models often struggle with real-time adaptability, which leads to inefficient decision-making and workflows. This paper will explore how autonomous AI agents enhance financial intelligence by leveraging reinforcement learning, multi-agent collaboration, and real-time behavioral analytics to optimize market predictions and risk management. Furthermore, regulatory compliance remains a major challenge, as opaque AI-driven models raise concerns about transparency, fairness, and ethical accountability. In this paper, we investigate the role of explainable AI (XAI), AI governance frameworks and bias-mitigation techniques, to ensure responsible deployment of AI in finance. Additionally, we will also discuss the operational inefficiencies in automated lending, wealth management, and algorithmic trading which are addressed through adaptive AI agents, capable of handling dynamic financial environments with minimal human intervention.

The hazards of AI-induced market disruptions, like flash crashes and systemic biases, are also examined in this research. To address these concerns, risk-aware AI architectures and real-time anomaly detection techniques are suggested.AI-powered financial assistants, which are intended for real-time data fusion, personalized wealth management, and intelligent portfolio rebalancing, are also used to address the lack of personalization in financial advice services. In order to ensure efficiency, compliance, security, and innovation in AI-driven financial services, this paper offers a thorough technological framework for utilizing the potential of agentic AI in financial ecosystems by addressing these possibilities and obstacles.

***Index Terms* - Agentic AI, Reinforcement Learning, RiskAware AI, Explainable AI (XAI), Financial Automation, AI Governance, Algorithmic Trading, AI-Powered Financial Assistants, Fraud Detection, Adaptive AI Agents**

I. INTRODUCTION

The finance industry has been a pioneer in adopting artificial intelligence (AI) to identify and mitigate risks, optimize decision-making, and improve operational efficiency. The traditional AI systems leveraging Machine Learning (ML), Deep Learning (DL), and Natural Language Processing (NLP), have already transformed financial services by enabling algorithmic trading, fraud detection, credit risk assessment, and roboadvisory platforms. However, these systems primarily function as decision-support tools, and require human oversight. As a result, these tools still struggle to adapt in real time to dynamic, stochastic financial environments [1]. It is with the emergence of Agentic AI which is characterized by autonomous decision making and self-learning capabilities that there has been a paradigm shift, allowing financial models to operate with minimal human intervention while continuously optimizing strategies based on evolving data patterns [1].

At the core of Agentic AI are advanced methodologies such as reinforcement learning (RL), deep reinforcement learning (DRL), multi-agent systems (MAS), and self-supervised learning (SSL), enabling financial systems to autonomously assess market conditions, execute trades, detect fraud, and optimize risk exposure. Unlike conventional AI models that rely on static training datasets and predefined rules, Agentic AI adjusts its decision-making framework in real-time using techniques like neural-symbolic learning, Bayesian inference, and Monte Carlo tree search (MCTS) [2]. In trading environments, AI driven hedge funds leverage these capabilities to refine portfolio allocations, while decentralized fraud detection mechanisms employ graph neural networks (GNNs) and generative adversarial networks (GANs) to proactively identify suspicious transactions. Additionally, regulatory compliance—a critical aspect of financial AI adoption—necessitates the integration of explainable AI (XAI) techniques, such as Shapley Additive Explanations (SHAP), Local Interpretable Model-agnostic Explanations (LIME), and counterfactual reasoning, to ensure transparency in AI-driven lending models, credit scoring, and automated risk assessments [3].

Despite its transformative potential, the deployment of Agentic AI in finance presents substantial challenges. The black-box nature of deep learning architectures, particularly in transformer-based financial models like FinBERT and GPT based trading systems, raises concerns about model interpretability and auditability [3]. Additionally, AI-driven financial decisions risk perpetuating systemic biases present in historical data, leading to unfair loan approvals, biased credit risk assessments, and

exclusionary investment strategies. Ethical concerns aside, regulatory and governance challenges remain a formidable barrier, as financial oversight bodies — including those shaping the EU AI Act, SEC AI Regulations, and ISO/IEC 42001 AI Risk Management standards—grapple with defining a clear framework for autonomous AI decision making in financial markets [4].

This paper explores the opportunities and challenges of Agentic AI in finance, examining its technical capabilities, applications, and regulatory implications. By analyzing real world use cases and industry adoption trends, we assess how financial institutions can harness self-learning financial agents while addressing risks related to explainability, bias, adversarial attacks, and compliance. Furthermore, we propose a future roadmap for responsible AI deployment, emphasizing the importance of XAI-driven financial models, ethical AI auditing frameworks, and adaptive regulatory policies to ensure a secure and efficient transition toward AI-driven finance.

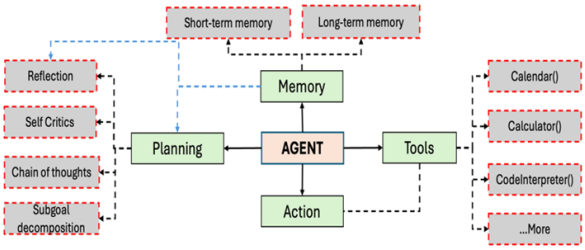

Figure 1:AI agent decision making

II. OPPORTUNITES OF AGENTIC AI IN FINANCE

Agentic AI represents a significant shift in financial technology, moving beyond assistive decision support to autonomous, adaptive, and proactive decision-making. Conventional AI primarily assists with data-driven recommendations but Agentic AI leverages deep reinforcement learning (DRL), multiagent systems (MAS), transformer-based language models (LLMs)and self-supervised learning (SSL) to assess, predict, and execute financial operations in real time. Its integration into finance has created groundbreaking opportunities across multiple domains, categorized below.

2.1 Autonomous Decision-Making in Trading and Risk Hedging

Agentic AI improves Algorithmic and High Frequency Trading (HFT) through self-learning agents that can optimize execution strategies in real-time. It uses Deep reinforcement learning (DRL) techniques like PPO and Q-learning that allow AI-driven hedge funds to dynamically adjust trading strategies based on market shifts [5]. Transformer-based models are used, such as FinBERT and GPT-Trader, which can enhance predictive analytics, improving market trend forecasting, risk adjusted returns, and arbitrage opportunities [6].

Use of multi-agent reinforcement learning (MARL) models can enable AI-driven portfolio management systems that analyze macroeconomic indicators, geopolitical risks, and sentiment data to adjust asset allocations dynamically. Bayesian inference models can predict black swan events and prevent excessive drawdowns during market crashes [1]. Agentic AI powered digital wealth advisors can revolutionize risk-aware investment strategies by continuously adapting to customer behaviors and financial goals [7].

2.2 Fraud Detection and Anti-Money Laundering (AML) Compliance

Traditional fraud detection relies on rule-based approaches which can fail to detect emerging fraud patterns. Agentic AI deploys graph neural networks (GNNs), anomaly detection models, and adversarial AI to autonomously identify fraudulent transactions in real-time. Multi-agent fraud detection models continuously cross-verify transactional anomalies, reducing false positives and improving fraud detection accuracy [5].
Agentic AI integrates behavioral analytics, entity resolution models, and self-supervised learning to detect unusual transaction patterns linked to money laundering. By leveraging explainable AI (XAI) methods like SHAP and LIME, AI driven AML systems ensure transparent compliance reporting while minimizing regulatory violations [8].

2.3 Personalized Customer Experience

Unlike traditional chatbot-advisors, Agentic AI-driven financial assistants leverage conversational AI, goal-driven reinforcement learning, and multimodal LLMs to offer context aware financial recommendations. These systems continuously learn and evolve in their response to user behaviors, optimizing financial planning without direct human intervention [1].

Agentic AI disrupts credit underwriting by incorporating alternative financial footprints like utility payments, e-commerce transactions, employment history into dynamic risk models. This shift improves financial inclusion, particularly for underbanked populations, while ensuring regulatory fairness through XAI-based bias monitoring [3].

2.4 Regulatory Compliance and Risk Governance

Agentic AI automates continuous monitoring of Regulations such as Basel III, GDPR, and ISO/IEC 42001 by integrating legal NLP models, predictive analytics, and anomaly detection frameworks to detect non-compliant transactions and regulatory risks before they escalate. Citigroup's report highlights how "Agentic AI auditors" will autonomously scan financial reports and contractual obligations, ensuring real-time compliance assessments [7].

Agentic AI can leverage neural-symbolic reasoning and multi-agent anomaly detection to autonomously scan transaction logs, financial statements, and corporate disclosures to identify fraudulent activities and accounting discrepancies. These AI models assist in detecting regulatory violations under IFRS, SOX, and MiFID II compliance frameworks [3].

2.5 Market Sentiment Analysis and Predictive Analytics

Agentic AI leverages integration with LLMs like FinBERT, and domain-specific NLP models to assess market sentiment from earnings reports, financial news, and social media trends. By leveraging contextual embeddings and reinforcement learning, AI models execute sentiment-driven trades with higher precision [6].

Financial crises, inflation trends, and systemic risks can be predicted by analyzing macroeconomic indicators and cross border trade data using Long short-term memory (LSTM) models, coupled with Transformer-based time-series forecasters. Additionally, AI-based Central bank models are emerging as a tool for monetary policy simulations and macro-financial stability analysis [3].

2.6 AI-Augmented Decision Support and Strategic Advisory

Unlike traditional chatbot-advisors, Agentic AI-driven financial assistants leverage conversational AI, goal-driven reinforcement learning, and multimodal LLMs to offer context aware financial recommendations. These systems continuously learn and evolve in their response to user behaviors, optimizing financial planning without direct human intervention [1].

Agentic AI disrupts credit underwriting by incorporating alternative financial footprints like utility payments, e-commerce transactions, employment history into dynamic risk models. This shift improves financial inclusion, particularly for underbanked populations, while ensuring regulatory fairness through XAI-based bias monitoring [3].

III. CHALLENGES OF USING AGENTIC AI IN FINANCE

Agentic AI presents a strong potential for transforming the finance industry, ranging from automating complex decision making to enhancing risk management processes. However, these advancements come with inherent challenges that need to be carefully addressed. These challenges span across data quality, interpretability, regulatory, compliance, operational integration and ethical decision-making standards.

3.1 Data Integrity and Quality

The diversity and incompleteness of financial data pose significant challenges for effective implementation of Agentic AI. Financial datasets consist of both structured (e.g., transaction logs, balance sheets) and unstructured data (e.g., social media sentiment, news feeds). Inconsistent formats, missing values, and differing structures across data sources can compromise the overall reliability of AI models [1]. For instance, credit risk models that rely on historical customer data can suffer from gaps in data, when integrating external sources like social media sentiment or unstructured market news, which affects their predictive accuracy. Additionally, imbalanced data distributions can lead to skewed predictions, where rare but critical financial events, such as fraud or market crashes, may be underrepresented. To mitigate these issues, techniques like Synthetic Minority Over-sampling Technique (SMOTE) or advanced imputation methods are applied to balance datasets. However, these solutions often increase computational demands, adding to the operational complexity [3].
In high-frequency trading (HFT), even microseconds of latency can result in massive financial losses. The computational intensity of deep learning algorithms, combined with the need to process vast streams of real-time data, creates a major bottleneck. Real-time processing requires also specialized hardware, such as highly computational GPUs, and well optimized algorithms to ensure that trading decisions are made within the time frame allowed [7]. If an AI model introduces significant latency in its decision-making process, the financial institution may miss trading opportunities or respond too late to emerging market threats.

Integrating heterogeneous data sources, such as trading data, market sentiment, and geopolitical news, into a unified system without introducing delays is another ongoing challenge. High-frequency models must process these inputs with minimal overhead, making real-time data engineering a crucial focus [3].

3.2 Algorithmic Bias and Model Interpretability

Algorithmic fairness in AI is a key challenge when implementing Agentic AI. The autonomous AI models used in credit scoring or loan approval systems, are at a risk of perpetuating biases inherent in historical data. A Forbes report mentioned that an AI Bias caused 80 percent of Black mortgage applicants to be rejected. Addressing these biases needs advanced fairness-enhancing techniques, such as adversarial debiasing or fairness constraints in the model optimization process.[9]

However, reduction of bias is not a one-size-fits-all process. For instance, enforcing fairness in a model which predicts creditworthiness may conflict with the model's objective of maximizing predictive accuracy, requiring a delicate trade-off between fairness and performance. Additionally, bias can stem from multiple sources, including the model design, data collection process, or even the features used in the model, necessitating a continuous cycle of audits and retraining [5].

Another critical challenge is the opacity of many AI models, especially in deep learning or reinforcement learning architectures. For financial institutions, it's crucial not only to deploy effective AI models but also to ensure that the decision-making process is explainable and transparent to stakeholders, including regulators and customers. Models such as neural networks or ensemble methods are often considered "black box" models because their decision-making process is not immediately understandable by humans (Towards Self-regulating AI, 2020 [5]).

In regulated sectors like finance, explainability is critical. Techniques like Shapley Additive Explanations (SHAP) and Local Interpretable Model-agnostic Explanations (LIME) have been proposed to provide transparency into how models make decisions. These techniques break down model outputs into more interpretable components, such as feature contributions, and are essential for ensuring that AI-based financial systems are auditable and accountable. However, these methods themselves add complexity and can sometimes be computationally expensive, making their widespread use a challenge [3].

3.3 Regulatory and Compliance Challenges

The regulatory environment surrounding AI in finance is still growing, and institutions need to remain agile in adhering to these emerging compliance standards. Strict rules on data protection, the explainability of AI judgments, and model audits are enforced by laws such as the General Data Protection Regulation (GDPR) in Europe and other regional finance regulations. The dynamic nature of AI means that financial institutions must be prepared for sudden regulatory changes that could impact their AI strategies [5].

AI-driven models particularly need to integrate explainability techniques which align with regulatory expectations. For example, the AI Act proposed by the European Union is likely to impose tighter scrutiny on AI applications in high-risk sectors like finance, requiring more stringent model explainability and governance [5]. To stay compliant, financial institutions must develop adaptive frameworks that allow for model retraining and explanation while ensuring that regulatory requirements are met at each stage of AI deployment.

Financial institutions may exploit differences in regulatory frameworks between jurisdictions, creating opportunities for regulatory arbitrage. This can lead to AI systems being deployed in one region that are not fully compliant with regulations in another [7]. Due to the worldwide interconnectedness of financial markets, this disparity may lead to instability if some market players take advantage of regulation gaps. To reduce arbitrage risks, financial organizations must adopt a global compliance approach in addition to ensuring compliance within their home jurisdiction [5].

3.4 Operational and System Integration

Many financial institutions still rely on legacy systems that were not designed for the integration with modern AI technologies. These systems often lack the necessary computational power or compatibility to support real-time AI decision-making. Integration support becomes further complicated when AI driven systems need to communicate with these legacy infrastructures to access critical data, like account information or market data feeds. Overcoming this challenge requires a significant investment in system upgrades, as well as the adoption of hybrid approaches where AI systems can work alongside traditional tools without disrupting ongoing operations [1]. To integrate AI into legacy systems, financial institutions are increasingly turning to microservices architecture, which decouples AI applications from core systems, enabling them to function independently while still interacting with traditional systems. However, this approach also introduces challenges related to data synchronization, security, and system latency [5].
The over-reliance on completely automated decision-making systems is a significant worry. Automation lowers human error and boosts efficiency, but it also runs the risk of removing human judgment from important financial decisions. For instance, an AI system may lack the nuanced understanding of a human expert in volatile market conditions or complex risk scenarios leading to suboptimal decisions [3]. Financial institutions need to strike a balance between autonomous decisions and human oversight, especially in situations where context, intuition, and expertise play a crucial role in decision making. Use of Hybrid AI systems that combine human oversight with algorithmic decision-making is becoming more common. Within these systems, AI models handle repetitive, data-heavy tasks, while humans provide critical oversight during exceptional situations [7].

IV. METHODOLOGIES USED IN AGENTIC AI FOR FINANCE

Agentic AI models require a robust training and fine-tuning approach to align with the complex and high-risk nature of the financial ecosystem. These methodologies must incorporate domain-specific learning, adversarial robustness, self-supervised techniques, and reinforcement learning to enhance AI's decision-making capabilities while ensuring compliance and transparency. This section explores how AI Agents can be trained for various financial applications, to ensure adaptability to dynamic market conditions and regulatory landscapes.

4.1 Supervised and Self-Supervised Learning for Credit Risk and Compliance

One of the basic methods for training AI models in finance is supervised learning, where models are trained using historical data with labeled outcomes. Credit risk assessment, for instance, benefits from models such as Gradient Boosted Decision Trees (GBDTs) and Deep Neural Networks (DNNs), trained on datasets like FICO, TransUnion, and Equifax. These models refine default probability predictions and loan approval decisions by identifying complex patterns in borrower profiles.[5]

Self-supervised learning (SSL) approach goes beyond traditional supervision. It is becoming increasingly important, especially in compliance and regulatory audits. SEC filings, Basel III reports, and MiFID II regulations are just a few of the large compliance datasets that AI models that have been refined using Masked Language Modeling (MLM) and contrastive learning may examine. For example, a study by Xinyu Li (2024) explores the integration of ML in credit risk analysis, which highlights how expansive language models like AntFinGLM, which are trained on financial texts and transactional data, can discern intricate financial patterns and market nuances, and this helps them in improving predictive accuracy [10].

4.2 Reinforcement Learning-Based Training for Algorithmic Trading

Reinforcement learning (RL) plays a pivotal role in training AI models for high-frequency trading and algorithmic market-making. Advanced RL methodologies, such as Inverse Reinforcement Learning (IRL), allow AI to infer optimal trading strategies by observing historical trader behavior. This approach helps AI anticipate market movements and adjust execution strategies dynamically, reducing slippage and improving profitability in volatile conditions [1]. In addition, Meta Reinforcement Learning (Meta-RL) enables AI trading models to generalize across multiple market conditions. Using Model Agnostic Meta-Learning (MAML), these models learn from diverse trading environments, allowing them to adapt rapidly to economic shifts such as recessions or financial crashes. This adaptability is critical for hedge funds and proprietary trading firms that rely on AI-driven execution engines [3].

4.3 Multi-Agent Learning for Market Simulations and Fraud Detection

Financial markets are inherently multi-agent environments where banks, hedge funds, and individual traders interact dynamically. Hierarchical Multi-Agent Reinforcement Learning (H-MARL) structures AI models into different layers, where high-level agents define strategic objectives while low-level agents optimize execution tactics. This approach enhances risk management, price discovery, and market stability, making AI-driven trading systems more resilient to unpredictable fluctuations [7].

Fraud detection in financial institutions is becoming intelligent from Federated Learning (FL), which allows multiple organizations to collaboratively train AI models while preserving data privacy. Federated AI models, trained using Secure Multiparty Computation (SMPC) and Differential Privacy, can detect anomalies in transaction data across multiple banks without sharing sensitive customer information. This methodology ensures compliance with privacy regulations such as GDPR and CCPA, making AI-driven fraud detection both effective and legally compliant [5].

4.4 Transfer Learning and Fine-Tuning LLMs for Financial Intelligence

Large Language Models (LLMs) are increasingly being fine-tuned for financial applications through transfer learning, where a pretrained model is adapted to a specific financial domain. For example, Bloomberg GPT and FinBERT are trained on earning reports, financial news, and investor sentiment datasets, enabling them to interpret market trends and macroeconomic indicators with domain-specific linguistic understanding [5].

Reinforcement Learning with Human Feedback (RLHF) is being employed to enhance AI-driven financial advisories. By continuously integrating expert feedback into AI's learning process, the AI model improves its investment recommendations while ensuring that fiduciary responsibilities and client specific risk profiles are maintained. RLHF-driven AI advisory systems optimize reward functions to balance profitability, risk management, and long-term financial planning.

The fine-tuning of Agentic AI models in finance requires a blend of supervised learning, reinforcement learning, multiagent collaboration, and transfer learning methodologies. These approaches ensure that AI can operate autonomously, adapt to evolving market conditions, comply with stringent regulations, and enhance decision-making across trading, compliance, fraud detection, and financial advisory services. As AI methodologies continue to advance, the integration of federated learning and human-aligned reinforcement learning will be critical in shaping the next generation of intelligent financial systems.

V. APPLICATIONS OF AGENTIC AI IN FINANCE

Agentic AI models require a robust training and fine-tuning approach to align with the complex and high-risk nature of the financial ecosystem. These methodologies must incorporate domain-specific learning, adversarial robustness, self-supervised techniques, and reinforcement learning to enhance AI's decision-making capabilities while ensuring compliance and transparency. This section explores how AI Agents can be trained for various financial applications, to ensure adaptability to dynamic market conditions and regulatory landscapes.

5.1 Algorithmic Trading and Market Making

Agentic AI plays a significant role in high-frequency trading (HFT), market making, and portfolio optimization by leveraging reinforcement learning (RL), deep learning, and multiagent systems to execute trades at optimal prices. Advanced AI models utilize Deep Q-Networks (DQNs), Proximal Policy

Optimization (PPO), and Transformer-based architectures to predict price movements, adjust trading strategies dynamically, and mitigate risks [3]. Agentic AI optimizes bid-ask spreads in market making by analyzing liquidity levels and adjusting positions autonomously. Techniques such as Multiagent Reinforcement Learning (MARL) allow multiple AI agents to simulate competitive market dynamics, improving the execution efficiency and liquidity provisioning [1].

5.2 Risk Management and Predictive Analytics

Financial institutions use Agentic AI for real-time risk assessment, stress testing, and anomaly detection by integrating self-learning models, adversarial robustness techniques, and explainable AI (XAI) frameworks. Advanced AI models like Graph Neural Networks (GNNs) can learn patterns that signify systemic risks[11] while Bayesian Inference and Monte Carlo Simulations allows for robust modeling of uncertainty in financial risk estimations[12].In credit risk management, Agentic AI-powered credit scoring models use Federated Learning (FL) and Differential Privacy (DP) to securely train on decentralized financial datasets, ensuring data privacy and regulatory compliance[7].

5.3 Fraud Detection and Financial Crime Prevention

Agentic AI enhances anti-money laundering (AML) compliance, fraud detection, and cybersecurity by employing unsupervised anomaly detection, graph-based fraud analytics, and self-supervised learning (SSL) models. Techniques such as Transformer-based anomaly detection, Autoencoders, and Generative Adversarial Networks (GANs) allow AI agents to detect fraudulent transactions in real-time [5].

Agent-based simulations help financial firms identify emerging fraud patterns by replicating fraudulent behaviors within synthetic environments. These simulations improve fraud detection models by adapting adversarial training strategies to counter evolving cyber threats [4].

5.4 Personalized Financial Advisory and Robo-Advisors

AI-driven robo-advisors leverage Agentic AI architectures to deliver personalized portfolio recommendations, retirement planning, and wealth management strategies. These AI systems incorporate Natural Language Processing (NLP) and Knowledge Graphs to understand client preferences, analyze financial documents, and provide customized investment strategies [1].

Autonomous financial advisors also use causal inference models to predict the impact of macroeconomic shifts, interest rate fluctuations, and market volatility on individual portfolios [13]. Multi-agent collaboration frameworks enable these AI-driven advisors to interact with hedge funds, brokerage systems, and financial analysts to optimize investment outcomes.

5.5 Regulatory Compliance and AI Governance

Agentic AI is transforming regulatory compliance by implementing automated regulatory reporting, compliance monitoring, and risk governance mechanisms. Explainable AI (XAI) techniques such as SHAP (Shapley Additive Explanations), LIME (Local Interpretable Model-Agnostic Explanations), and Counterfactual Reasoning ensure transparency in AI-driven financial decisions [3].

Financial institutions use RegTech AI systems to automate regulatory reporting, monitor trading activities, and identify compliance risks using Neural-Symbolic AI and Knowledge Graphs. These models scan regulatory updates in real time, ensuring that financial organizations adapt their AI systems to evolving compliance frameworks [7].

5.6 Loan Underwriting and Credit Scoring

AI-driven loan underwriting models use multi-modal learning to analyze creditworthiness from structured and unstructured data sources, including bank transactions, income patterns, and digital footprints. Hybrid AI architectures combining Recurrent Neural Networks (RNNs) and Graph Neural Networks (GNNs) provide a more comprehensive credit risk assessment [11].

Self-adaptive underwriting AI models leverage Bayesian Optimization and Transfer Learning to fine-tune decision making based on economic shifts and borrower behavior patterns. Additionally, privacy-preserving techniques such as Secure Multi-Party Computation (SMPC) allow financial institutions to collaborate on risk analysis without exposing sensitive customer data [5].

5.7 Decentralized Finance (DeFi) and Smart Contracts

Agentic AI is powering DeFi ecosystems by enabling intelligent automated market makers (AMMs), AI-optimized staking mechanisms, and AI-driven risk assessment models for blockchain-based financial platforms. Smart contract execution benefits from Reinforcement Learning-based contract optimization, ensuring that DeFi lending, staking, and liquidity pools operate efficiently [4].

Additionally, AI-driven DeFi risk models use on-chain and off-chain data fusion techniques to monitor smart contract vulnerabilities, liquidity risks, and token price manipulation. These AI models leverage Federated Learning and Zero Knowledge Proofs (ZKP) to ensure secure and privacy preserving financial transactions [7].

VI. CONCLUSION

The financial sector is undergoing a paradigm shift with the rise of Agentic AI, which is revolutionizing trading strategies, risk assessment, fraud detection, and customer interaction. By leveraging reinforcement learning, multi-agent collaboration, and federated learning, financial institutions can achieve unprecedented levels of efficiency, accuracy, and automation. However, this technological leap also introduces significant challenges, including AI explainability, ethical responsibility, regulatory compliance, and systemic risks associated with AI driven market volatility.

This paper has demonstrated that the key to safeguarding financial AI adoption lies in robust governance mechanisms, continuous fine-tuning through human oversight, and the implementation of privacy-preserving AI architectures. The deployment of self-regulating Agentic AI models, combined with explainable AI frameworks and reinforcement learning strategies, can ensure both efficiency and accountability in AI driven financial operations.

As financial markets continue to embrace Agentic AI, a multidisciplinary approach is required—one that blends technical innovation with ethical responsibility and regulatory foresight. Future research must focus on improving AI adaptability through self-supervised and adversarial learning techniques, ensuring that these autonomous systems remain resilient to emerging financial threats while maintaining compliance with global regulations. The evolution of Agentic AI in finance will not only

reshape economic infrastructures but also challenge traditional paradigms of financial governance, necessitating a holistic and proactive approach toward AI-driven financial transformation.

REFERENCES

[1] G. Samdani, Y. Dixit, and G. Viswanathan, "Agentic ai in autonomous financial advisories," World Journal of Advanced Engineering Technology and Sciences, vol. 9, no. 1, pp. 410–420, June 2023

[2] Z. Liang, F. Wei, W. Xu, L. Chen, Y. Qian, and X. Wu, "I-MCTS: Enhancing Agentic AutoML via Introspective Monte Carlo Tree Search," 2025. [Online]. Available: https://arxiv.org/abs/2502.14693

[3] C. Maple, L. Szpruch, G. Epiphaniou, K. Staykova, S. Singh, W. Penwarden, Y. Wen, Z. Wang, J. Hariharan, and P. Avramovic, "The AI Revolution: Opportunities and Challenges for the Finance Sector," 2023. [Online]. Available: https://arxiv.org/abs/2308.16538

[4] E. Kurshan, H. Shen, and J. Chen, "Towards self-regulating AI: challenges and opportunities of AI model governance in financial services," in Proceedings of the First ACM International Conference on AI in Finance, ser. ICAIF '20. New York, NY, USA: Association for Computing Machinery, 2021. [Online]. Available: https://doi.org/10.1145/3383455.3422564

[5] I. Okpala, A. Golgoon, and A. R. Kannan, "Agentic AI Systems Applied to tasks in Financial Services: Modeling and model risk management crews," 2025. [Online]. Available: https://arxiv.org/abs/2502.05439

[6] O. Shobayo, S. Adeyemi-Longe, O. Popoola, and B. Ogunleye, "Innovative Sentiment Analysis and Prediction of Stock Price Using FinBERT, GPT-4 and Logistic Regression: A Data-Driven Approach," Big Data and Cognitive Computing, vol. 8, no. 11, p. 143, Oct. 2024. [Online]. Available: http://dx.doi.org/10.3390/bdcc8110143

[7] R. Ghose, S. Bantanidis, K. Master, R. S. Shah, S. Chahal, P. Sharma, and C. Zhai, "Agentic AI: Finance & the 'Do It For Me' Economy," Citi GPS: Global Perspectives & Solutions, Tech. Rep., January 2025.

[8] M. A. Rahman and M. A. Hossain, "A Comprehensive Survey on Machine Learning Approaches in IoT Security" IEEE Internet of Things Journal, vol. 9, no. 1, pp. 1–24, 2022. [Online]. Available: https://ieeexplore.ieee.org/document/10705194

[9] K. Hale, "AI Bias Caused 80% of Black Mortgage Applicants to Be Denied," Forbes, September 2021

[10] X. Li, "Integration of Machine Learning in Credit Risk Analysis: Enhancing Predictive Accuracy through Financial Text Pretraining," ACE Proceedings, 2024

[11] L. Gonon, T. Meyer-Brandis, and N. Weber, "Computing Systemic Risk Measures with Graph Neural Networks," 2024. [Online]. Available: https://arxiv.org/abs/2410.07222

[12] S. Takeda and T. Kitada, "Bayesian Inference Based on Monte Carlo Technique," ASME Journal of Risk and Uncertainty in Engineering Systems, Part B: Mechanical Engineering, vol. 10, no. 4, p. 041203, 2024.

[13] L. G. Neuberg, "CAUSALITY: MODELS, REASONING, AND INFERENCE, by Judea Pearl, Cambridge University Press, 2000" Econometric Theory, vol. 19, no. 4, p. 675–685, 2003